

Into to C/C++ pragma-based parallelism

Igor Sfiligoi – San Diego Supercomputing Center (UCSD)

We will mostly focus on C

- “Full” C++ hides a lot of details from the final user
 - In C, you explicitly tell how to do things, step by step
 - C better for understanding how things work
- Most of the presented concepts still apply to “full” C++
 - But you may need to understand the minute details that the compiler normally tries to hide from you!
 - Wherever appropriate, will call out known “gotchas”

We will focus on portable methods

- Life is too short to re-implement your apps for every architecture!
 - That's why we use a compiler and not ASM to start with!
- Picking a vendor specific toolkit may give you better performance
But with dominant vendors changing, hard to maintain/port
 - CPUs - pthreads
 - NVIDIA GPUs – CUDA
 - AMD GPUs – ROCm/HIP
 - Intel GPUs – OneAPI/SYCL

We will focus on portable methods

- Life is too short to re-implement our apps for every architecture!
- We will start with!

**No magic solutions, some
porting/tuning always
needed with new platforms.**

- better performance
- need to maintain/port
- NVIDIA GPUs – CUDA
- AMD GPUs – ROCm/HIP
- Intel GPUs – OneAPI/SYCL

We will focus on portable methods

- Life is too short to re-implement your apps for every architecture!
 - That's why we use a compiler and not ASM to start with!
- Picking a vendor specific solution is not portable
But with dominant vendors, we can focus on mature solutions, so you can start using them now.
(Some cutting-edge solutions do promise to make your life easier)
 - CPUs - pthreads
 - NVIDIA GPUs – CUDA
 - AMD GPUs – ROCm/HIP
 - Intel GPUs – OneAPI/SYCL

Virtually all parallelism in loops

- Compilers have hard time parallelizing sequential code
 - Not completely impossible, but very limited potential
 - Sometimes you can help
 - But often makes code hard to maintain
- Loops are the natural concept for expressing parallelism, if
 - Each iteration independent
 - Number of iterations known in advance

```
a1 += 2*b1;  
a2 += 2*b2;  
a3 += 2*b3;  
a4 += 2*b4;
```

```
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
}
```

Virtually all parallelism in loops

- Compilers have hard time parallelizing sequential code
 - Not completely impossible, but very limited potential
 - Sometimes you can help
 - But often makes code hard to maintain

```
a1 += 3*b1;  
a2 += 3*b2;  
a3 += 3*b3;  
a4 += 3*b4;
```

- Loops are the natural concept for expressing parallelism, if
 - Each iteration independent
 - Number of iterations known in advance

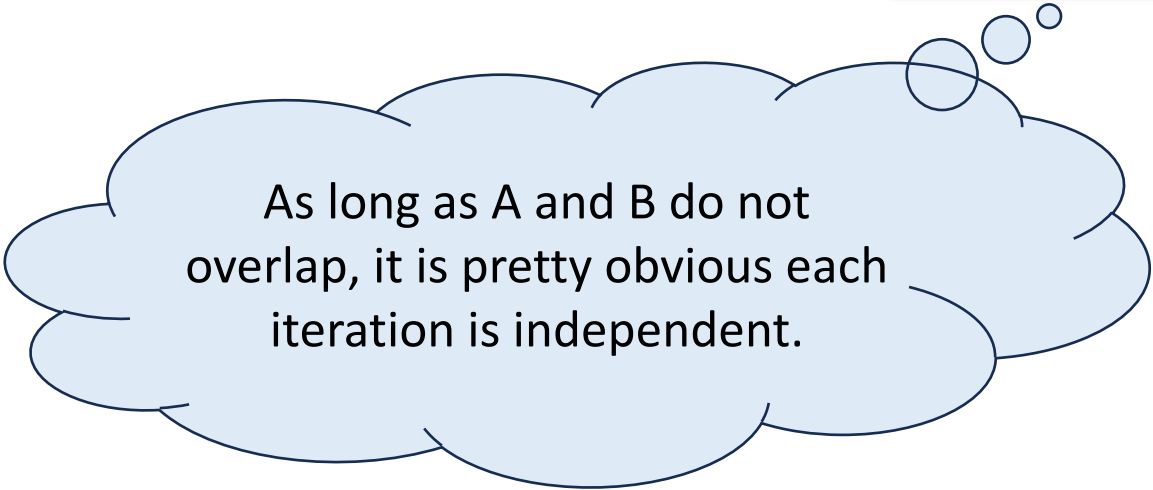
```
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
}
```

A good compiler will issue
vectorized code for both snippets

Loops and dependencies

- Loops are the natural concept for expressing parallelism, if
 - Each iteration independent
 - Number of iterations known in advance

```
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
}
```



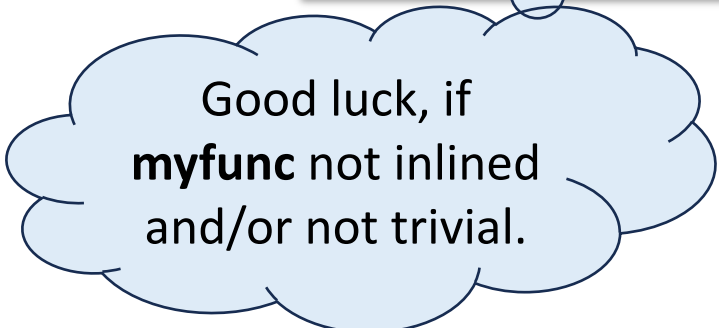
As long as A and B do not overlap, it is pretty obvious each iteration is independent.

Dependencies hard to infer

- Loops are the natural concept for expressing parallelism, if
 - Each iteration independent
 - Number of iterations known in advance
- Doesn't take much for compiler to give up
 - 1st priority for the compiler is to generate correct code
 - Optimization/speed is distant 2nd

```
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
}
```

```
for (int i=0; i<N; i++) {  
    A[i] += 3*myfunc(A,B,i)  
}
```



Good luck, if
myfunc not inlined
and/or not trivial.

Dependencies hard to infer

- Loops are the natural concept for expressing parallelism, if
 - Each iteration independent
 - Number of iterations known in advance
- Doesn't take much for compiler to give up
 - 1st priority for the compiler is to generate correct code
 - Optimization/speed is distant 2nd
- Not to mention “smart” users

```
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
}
```

```
for (int i=0; i<N; i++) {  
    A[i] += 3*myfunc(A,B,i)  
}
```

```
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
    if(A[i]>3) N=A[i];  
}
```

Helping the compiler with pragmas

- OpenMP has emerged as the major pragma-based parallelization paradigm
 - Especially on CPUs
 - Just decorate the loop with a pragma
- Explicitly tell the compiler that parallelization is both possible and desirable



<https://www.openmp.org>

```
#pragma omp parallel for
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
}
```

```
#pragma omp parallel for
for (int i=0; i<N; i++) {
    A[i] += 3*myfunc(A,B,i)
}
```

Helping the compiler with pragmas



<https://www.openmp.org>

- OpenMP is a portable API for parallel programming
- Just add pragmas to your code
- Explicitly parallelizing code is both possible and desirable

Compiler will trust you!
If the code was not independent,
you will get unexpected results.

```
#pragma omp parallel for
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
}
```

```
#pragma omp parallel for
for (int i=0; i<N; i++) {
    A[i] += 3*myfunc(A,B,i)
}
```

GPU extensions

- OpenMP recently added GPU support
 - The notion of “OMP TARGET”
- Slightly different syntax, but similar in concept

```
#pragma omp target teams distribute parallel for simd
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
}
```



<https://www.openmp.org>

GPU pragmas - OpenACC

- OpenACC was the pioneer in GPU-based pragma-based parallelization
- Only supported by NVIDIA compilers

```
#pragma acc parallel loop gang vector  
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
}
```



<https://www.openacc.org>

Very unfortunate.

Much easier to use than
OpenMP Target in my opinion.

Remember the launch cost

- Every parallel section incurs a “launch cost”
 - Very expensive on GPUs
 - But not negligible on CPUs, either
- Pack as much as possible in a single parallel section
 - Can be order of magnitude faster

```
#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
}

#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    C[i] += 4*D[i]
}
```



```
#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
    C[i] += 4*D[i]
}
```

But don't break semantics!

- Every parallel section incurs a “launch cost”
 - Very expensive on GPUs
 - But not negligible on CPUs, either
- Pack as much as possible in a single parallel section
 - Can be order of magnitude faster
 - **But make sure you preserve the independence of the iterations!**

```
#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
}
#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    C[i] += 4*A[1+(2*i+5)%N]
}
```



```
#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
    C[i] += 4*A[1+(2*i+5)%N]
}
```


Remember memory locality

- Do not iterate over same buffer multiple times
 - If at all possible (dependencies)
- Pack as much as possible in a single parallel section
 - Also use temp variables when appropriate
 - **Will be order of magnitude faster on large buffers**

```
#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
}

#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    C[i] += 4*A[i] + 3*B(i)
}
```



```
#pragma omp target teams \
  distribute parallel for simd
for (int i=0; i<N; i++) {
    b3 = 3*B[i]
    a2 = A[i] + b3
    A[i] = a2
    C[i] += 4*a2 + b3
}
```

Dealing with partitioned memory

- Remainder: CPU and GPU memories are (typically) independent
- By default, OpenMP will copy all buffers
 - From CPU to GPU memory before GPU kernel start
 - From GPU to CPU memory after GPU kernel completion
- If buffers only needed on one side, you can avoid the transfer
 - Resulting in drastic speedup!

} Slow!

```
#pragma omp target enter data map(to:A)
...
#pragma omp target teams distribute parallel for simd \
    map(tofrom:A) map(to:B)
for (int i=0; i<N; i++) {
    A[i] += 3*B[i]
}
```

Dealing with partitioned memory

- OpenMP Target memory uses reference counters
 - Will only transfer once, even if mentioned several times
- 3+1 possible directions:
 - “to” – CPU -> GPU
 - “from” – GPU -> CPU
 - “tofrom” – both ways
 - May create uninitialized buffer

```
...
#pragma omp target enter data map(to:A)
...
#pragma omp target teams distribute parallel for simd \
    map(tofrom:A) map(to:B)
    for (int i=0; i<N; i++) {
        A[i] += 3*B[i]
    }
...
#pragma omp target exit data map(delete:A)
```

Dealing with partitioned memory

- OpenMP Target memory uses reference counters

- Will eventually use GPU's "A"

- 3+1 possible execution flows:

- "Copy "B" CPU->GPU"
- to

- May cr

Release GPU's "B"

Release GPU's "A"

```
...  
#pragma omp target enter data map(to:A)  
...  
#pragma omp target teams distribute parallel for simd \  
    map(tofrom:A) map(to:B)  
    for (int i=0; i<N; i++) {  
        A[i] += 3*B[i]  
    }  
...  
#pragma omp target exit data map(delete:A)
```

Dealing with partitioned memory

- OpenMP Target memory uses references

- Will eventually use GPU's "A"

- 3+1 possible actions:

- "Copy "B" CPU->GPU"
- to

- May cr

Release GPU's "B"

Release GPU's "A"

Copy "A" CPU->GPU

In this example,
CPU's "A" never updated

```
...  
#pragma omp target enter data map(to:A)  
...  
#pragma omp target teams distribute parallel for simd \  
map(tofrom:A) map(to:B)  
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
}  
...  
#pragma omp target exit data map(delete:A)
```

Managed memory

- Remainder: CPU and GPU memories are (typically) independent
 - **But most modern GPUs can hide this through managed memory**
 - **No need to explicitly manage memory in that case**
- Only works on buffers explicitly allocated as managed
 - Either explicitly,
by replacing malloc with e.g. (not portable)
`cudaMallocManaged`
 - Or **implicitly**,
by using compiler directives, e.g.
`nvc++ -gpu=managed`
- GPU driver will page those buffers as needed between CPU and GPU memories

```
cudaMallocManaged(&A,N);  
...  
#pragma omp target teams distribute parallel for simd  
for (int i=0; i<N; i++) {  
    A[i] += 3*B[i]  
}
```

Managed memory

- Remainder: CPU and GPU memories are (typically) independent
 - **But most modern GPUs can hide this through managed memory**
 - **No need to explicitly manage memory in that case**
- Only works on buffers explicitly allocated as managed
 - Either explicitly,
by replacing malloc with e.g. (not portable)
`cudaMallocManaged`
 - Or **implicitly**,
by using compiler directives, e.g.
`nvc++ -gpu=managed`
- GPU driver will page those buffers as needed between CPU and GPU memories



Be careful not to
start trashing!
(swap like)

```
cudaMallocManaged  
...  
#pragma omp parallel for  
for (int i=0; i<N; i++)  
    A[i] += 3*B[i]  
}
```

Ready for some hands-on?